



Fatih Yegen

Electrics And Electronics Software Engineer

Aydın, Turkey • fatihyegen.09@gmail.com • 0505 034 00 09 • fygen.github.io • fatihyegen
fygen

About

Each section title is arbitrary.

You can choose any of the 9 entry types for each section.

Markdown syntax is supported everywhere. This is **bold**, *italic*, and link.

Education

- BS 42 Istanbul**, BS in Computer Science – Istanbul, TR Nov 2023 – Nov 2025
Embedded engineering OOP, C, C++ , Bash, Docker, Network engineering, terminal Open Source projects and others.
 - Half automation of one school project. Born2BeRoot.
 - Understanding of Tech Things
- BS Pamukkale University**, BS in Electrics & Electronics Engineering – Denizli, Türkiye Sept 2014 – June 2019
Meeting with Electrics and Electronics.
 - GPA: 3.14/4.00
 - YDS: 66
 - The teacher of classes award.
 - Open sourced*
- Eras-mus Warsaw University Of Technology**, Erasmus in Electrical Engineering & Electronics and Information Technology – Warsaw, PL 2016 – 2017
 - GPA: 70/100
 - 2 semesters
 - Most people person award.

Experience

- Sürat Kargo**, Software Development Specialist – İstanbul, Türkiye June 2023 – Jan 2025
Software Developer Stack I am currently interested in this stack. Want to continue my journey with, LLM APIs (LLM APIs for DB control, query and so.), Robotics, anything including; improving our daily lives.
 - Built foundation model infrastructure serving 2M+ monthly API requests with 99.97% uptime
 - Raised \$18M Series A led by Sequoia Capital, with participation from a16z and Founders Fund
 - Scaled engineering team from 3 to 28 across ML research, platform, and applied AI divisions
 - Developed proprietary inference optimization reducing latency by 73% compared to baseline
- NVIDIA Research**, Research Intern – Santa Clara, CA May 2022 – Aug 2022
 - Designed sparse attention mechanism reducing transformer memory footprint by 4.2x
 - Co-authored paper accepted at NeurIPS 2022 (spotlight presentation, top 5% of submissions)
- Google DeepMind**, Research Intern – London, UK May 2021 – Aug 2021
 - Developed reinforcement learning algorithms for multi-agent coordination
 - Published research at top-tier venues with significant academic impact
 - ICML 2022 main conference paper, cited 340+ times within two years
 - NeurIPS 2022 workshop paper on emergent communication protocols
 - Invited journal extension in JMLR (2023)
- Apple ML Research**, Research Intern – Cupertino, CA May 2020 – Aug 2020
 - Created on-device neural network compression pipeline deployed across 50M+ devices
 - Filed 2 patents on efficient model quantization techniques for edge inference

Microsoft Research, Research Intern – Redmond, WA

May 2019 – Aug 2019

- Implemented novel self-supervised learning framework for low-resource language modeling
- Research integrated into Azure Cognitive Services, reducing training data requirements by 60%

Projects

FlashInfer

Jan 2023 – present

Open-source library for high-performance LLM inference kernels

- Achieved 2.8x speedup over baseline attention implementations on A100 GPUs
- Adopted by 3 major AI labs, 8,500+ GitHub stars, 200+ contributors

NeuralPrune

Jan 2021

Automated neural network pruning toolkit with differentiable masks

- Reduced model size by 90% with less than 1% accuracy degradation on ImageNet
- Featured in PyTorch ecosystem tools, 4,200+ GitHub stars

Publications

Sparse Mixture-of-Experts at Scale: Efficient Routing for Trillion-Parameter Models

July 2023

John Doe, Sarah Williams, David Park

10.1234/neurips.2023.1234 (NeurIPS 2023)

Neural Architecture Search via Differentiable Pruning

Dec 2022

James Liu, John Doe

10.1234/neurips.2022.5678 (NeurIPS 2022, Spotlight)

Multi-Agent Reinforcement Learning with Emergent Communication

July 2022

Maria Garcia, John Doe, Tom Anderson

10.1234/icml.2022.9012 (ICML 2022)

On-Device Model Compression via Learned Quantization

May 2021

John Doe, Kevin Wu

10.1234/iclr.2021.3456 (ICLR 2021, Best Paper Award)

Selected Honors

- MIT Technology Review 35 Under 35 Innovators (2024)
- Forbes 30 Under 30 in Enterprise Technology (2024)
- ACM Doctoral Dissertation Award Honorable Mention (2023)
- Google PhD Fellowship in Machine Learning (2020 – 2023)
- Fulbright Scholarship for Graduate Studies (2018)

Skills

Languages: Python, C++ , CUDA, Rust, Julia

ML Frameworks: PyTorch, JAX, TensorFlow, Triton, ONNX

Infrastructure: Kubernetes, Ray, distributed training, AWS, GCP

Research Areas: Neural architecture search, model compression, efficient inference, multi-agent RL

Patents

- Adaptive Quantization for Neural Network Inference on Edge Devices (US Patent 11,234,567)
- Dynamic Sparsity Patterns for Efficient Transformer Attention (US Patent 11,345,678)
- Hardware-Aware Neural Architecture Search Method (US Patent 11,456,789)

Invited Talks

- Scaling Laws for Efficient Inference — Stanford HAI Symposium (2024)
- Building AI Infrastructure for the Next Decade — TechCrunch Disrupt (2024)
- From Research to Production: Lessons in ML Systems — NeurIPS Workshop (2023)
- Efficient Deep Learning: A Practitioner’s Perspective — Google Tech Talk (2022)